

PROBLEM STATEMENT: Design a distributed application using MapReduce which processes a log file of a system. List out the users who have logged for maximum period on the system. Use simple log file from the Internet and process it using a pseudo distribution mode on Hadoop platform.

Step 1: Create a Java Project in Eclipse

1. Create a New Java Project:

- Open Eclipse and go to **File** → **New** → **Java Project**.
- Name the project (e.g., `HadoopLogFileProcessor`).
- Click **Finish**.

2. Create a Package:

- In the **Project Explorer**, right-click on `src` → **New** → **Package**.
- Name the package (e.g., `HadoopLogFileProcessor`).

3. Create the Classes:

- Right-click on the package you created → **New** → **Class**.
- Create the following classes:
 - **LoginMapper**
 - **LoginReducer**
 - **ProcessLogFile** (Driver class)

Step 2: Add the Code

Add the following code in their respective classes:

//LoginMapper.java

```
package HadoopLogFileProcessor;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Mapper;
```

```
import java.io.IOException;

public class LoginMapper extends Mapper<Object, Text, Text, IntWritable> {

    // For emitting user login and logout events

    private Text user = new Text();

    private IntWritable timestamp = new IntWritable();

    public void map(Object key, Text value, Context context) throws IOException,
    InterruptedException {

        String line = value.toString();

        String[] parts = line.split(" ");

        // Extract timestamp, user, and action

        int time = Integer.parseInt(parts[0]);

        String username = parts[1];

        String action = parts[2];

        if (action.equals("login")) {

            // Emit the login timestamp with key as <username>_login

            user.set(username + "_login");

            timestamp.set(time);

            context.write(user, timestamp);

        } else if (action.equals("logout")) {

            // Emit the logout timestamp with key as <username>_logout

            user.set(username + "_logout");

            timestamp.set(time);

            context.write(user, timestamp);

        }

    }

}
```

```
    }  
  }  
}
```

//LoginReducer.java

```
package HadoopLogFileProcessor;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Reducer;
```

```
import java.io.IOException;
```

```
import java.util.HashMap;
```

```
import java.util.Map;
```

```
public class LoginReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
```

```
    private Map<String, Integer> userSessionDurations = new HashMap<>();
```

```
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws  
    IOException, InterruptedException {
```

```
        String username = key.toString().split("_")[0];
```

```
        // Iterate through all timestamps for login/logout
```

```
        int loginTimestamp = -1;
```

```
        int logoutTimestamp = -1;
```

```
        for (IntWritable val : values) {
```

```
            if (key.toString().endsWith("login")) {
```

```

        loginTimestamp = val.get();
    } else if (key.toString().endsWith("logout")) {
        logoutTimestamp = val.get();
    }
}

```

```

// Calculate session duration if both login and logout are present
if (loginTimestamp != -1 && logoutTimestamp != -1) {
    int sessionDuration = logoutTimestamp - loginTimestamp;
    userSessionDurations.put(username, sessionDuration);
}
}

```

// Output the user with the longest session

@Override

protected void cleanup(Context context) throws IOException, InterruptedException {

```

    String maxUser = "";

```

```

    int maxDuration = 0;

```

```

    for (Map.Entry<String, Integer> entry : userSessionDurations.entrySet()) {

```

```

        if (entry.getValue() > maxDuration) {

```

```

            maxDuration = entry.getValue();

```

```

            maxUser = entry.getKey();

```

```

        }

```

```

    }

```

```

        context.write(new Text("User with longest session: " + maxUser), new
        IntWritable(maxDuration));

```

```
}  
}
```

//ProcessLogFile.java

```
package HadoopLogFileProcessor;  
  
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Job;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
  
public class ProcessLogFile {  
  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "Log File Processor");  
  
        // Set the Jar class  
        job.setJarByClass(ProcessLogFile.class);  
  
        // Set Mapper and Reducer classes  
        job.setMapperClass(LoginMapper.class);
```

```

    job.setReducerClass(LoginReducer.class);

    // Set output key and value types
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);


    // Set input and output paths
    FileInputFormat.addInputPath(job, new Path(args[1]));
    FileOutputFormat.setOutputPath(job, new Path(args[2]));


    // Wait for job completion
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Step 3: Build the Project

1. Build the Project:

- Click on **Project** → **Build Project** in Eclipse.
- Make sure the project compiles without any errors.

2. Check Build Path:

- Go to **Project Explorer** → Right-click on your project → **Build Path** → **Configure Build Path**.
- Ensure that all the Hadoop JAR files you added are present in the **Libraries** section.
 - i. `adoop-common` (e.g., `hadoop-common.jar`)
 - ii. `Hadoop-mapreduce-client-core` (e.g., `hadoop-mapreduce-client-core-2.x.x.jar`)

Step 4: Create the JAR File

1. Export to JAR:

- Go to **File** → **Export**.
- Choose **Java** → **Runnable JAR file**.
- Choose **Launch configuration** as **ProcessLogFile**.
- Select the destination path and name your JAR file (e.g., **Log.jar**).
- Click **Finish**.

Step 5: Prepare the Input Files

1. Create the Log File (**logfile.txt**):

- The **logfile.txt** should be placed in the HDFS directory. Here's the content of the file:

1624875600 user1 login

1624875700 user2 login

1624875800 user1 logout

1624875900 user2 logout

1624876000 user3 login

1624876100 user3 logout

Upload the Input File to HDFS:

- Use the Hadoop shell to upload the input file to HDFS:

```
hdfs dfs -put /path/to/logfile.txt /user/cloudera/logfile.txt
```

Step 6: Configure the Run Configuration

1. Set up Run Configuration:

- In Eclipse, go to **Run** → **Run Configurations**.
- Select **Java Application** and click **New**.

- In the **Main** tab, select the **Project** (your Hadoop project) and the **Main Class** (`ProcessLogFile`).

Step 8: Run the Job

1. Run the Hadoop Job:

```
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/Log.jar
```

```
HadoopLogFileProcessor.ProcessLogFile /user/cloudera/logfile.txt /user/cloudera/dir51
```

2. hdfs dfs -ls /user/cloudera/dir51

```
Bytes Written=30
[cloudera@quickstart ~]$ hdfs dfs -ls /user/cloudera/dir51

Found 2 items
-rw-r--r--  1 cloudera cloudera      0 2025-04-26 23:53 /user/cloudera/dir51/_SUCCESS
-rw-r--r--  1 cloudera cloudera    30 2025-04-26 23:53 /user/cloudera/dir51/part-r-00000
[cloudera@quickstart ~]$
```